

Highlights

- Imputation, feature selection and stacking were all fitted inside validation.
- Stacking did not beat logistic regression (CV AUC 0.843 vs 0.842).
- RFECV cut inputs by 25%, keeping a stable core of glucose and BMI terms.
- Probabilities were too high, consistent with class weighting.
- Sensitivity-targeted thresholds transferred from development to hold-out data.